

AWS Bedrock – deploy and consume a LLM

Go to AWS Console → Lambda → Create Function

- **Name:** bedrock-gateway
- **Runtime:** Python 3.10
- **Execution Role:** Create a new role with basic Lambda permissions

Paste following code:

```
import json
```

```
import boto3
```

```
def lambda_handler(event, context):
```

```
    body = json.loads(event['body'])
```

```
    prompt = body.get("prompt", "")
```

```
    client = boto3.client("bedrock-runtime", region_name="us-east-1")
```

```
    response = client.invoke_model(
```

```
        modelId="anthropic.claude-3-sonnet-20240229-v1:0",
```

```
        contentType="application/json",
```

```
        accept="application/json",
```

```
        body=json.dumps({
```

```
            "prompt": f"\n\nHuman: {prompt}\n\nAssistant:",
```

```
            "max_tokens_to_sample": 300,
```

```
            "temperature": 0.7
```

```
        })
```

```
    )
```

```
result = response['body'].read().decode("utf-8")

return {

    "statusCode": 200,

    "headers": {"Content-Type": "application/json"},

    "body": result

}
```

Test Lambda in Console

1. Go to **"Test" tab**
2. Configure a new test event (name: testPrompt)
3. Paste this as the test input:

```
{

    "body": "{\"prompt\": \"What is quantum computing?\"}"

}
```

4. Click **Test**

You should see 200 status and Claude's answer in output

Create API Gateway to Expose Lambda as REST API

Go to AWS Console → API Gateway → Create API

- Choose **HTTP API**
- Integration: **Lambda Function** → bedrock-gateway
- Route: POST /generate
- Deploy → Name the stage: prod

Copy the **Invoke URL**, e.g.:

<https://abc123xyz.execute-api.us-east-1.amazonaws.com/generate>

Test API Gateway from Console

1. Go to **API Gateway > Your API > Stages > prod**

2. Click **“Invoke URL”** → it will open in browser (it won’t work directly, it’s a POST API)

3. Go to **Postman** or use the **Test Client in Console**:

- In **API Gateway**, go to **Routes > /generate**
- Click **“Test”**
- Set Method: POST
- Headers:

Content-Type: application/json

Body:

```
{  
  "prompt": "Tell me about black holes."  
}
```

Click **Send Test Request**

● You should get a 200 response with the Claude model's output.